Task 3P

Step One: Creating a Managed Identity:

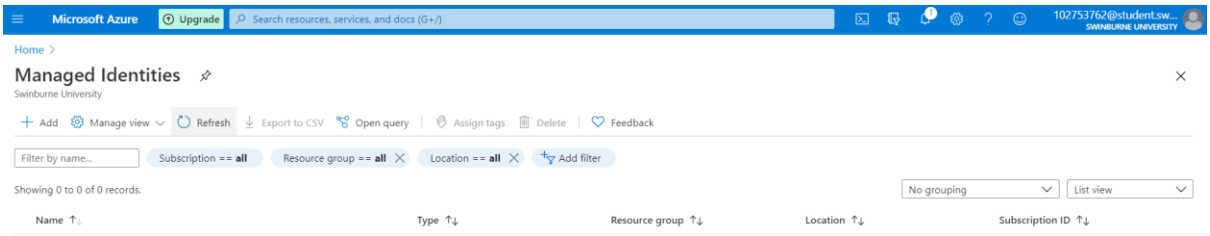

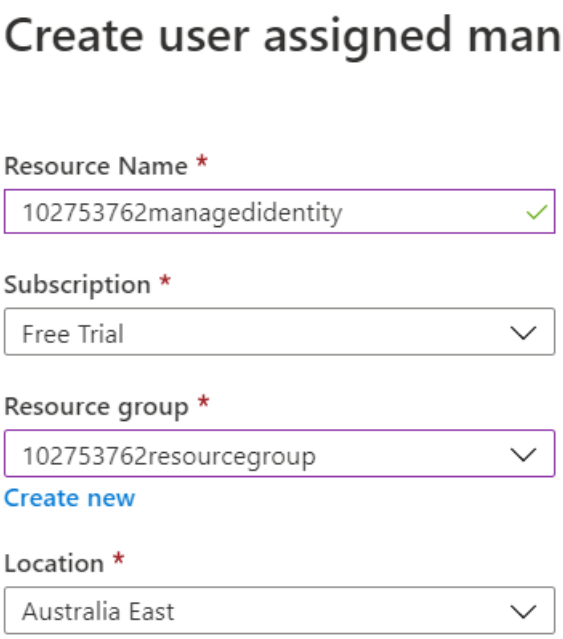

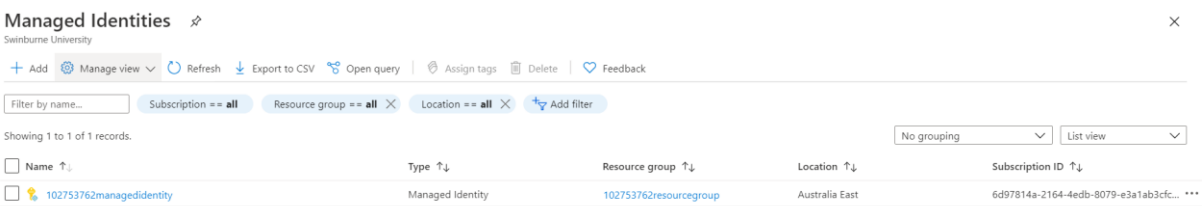

Question 1: What is a data warehouse, and how does a data lake differ from it?

Answer: According to IBM, “A data warehouse is a system that aggregates data from different sources into a single, central, consistent data store to support business analytics, data mining, artificial intelligence (AI), and machine learning”.

On the other hand, according to redhat, “A data lake is a type of data repository that stores large and varied sets of raw data in its native format. Data lakes let you keep an unrefined view of your data”.

Difference between data warehouse and data lake:

01. When developing data warehouse, data source is analysed and it is a structured data model. Not all data from the source is stored in the data warehouse. So only data that will be used to answer a specific question or will be included in a report are stored in data warehouse. On the other hand, all data are stored in data lake not just the data that will be used for analysing or reporting.
02. All type of data can be stored in data lake including non-traditional data like, web server log, sensor data, social network activity etc. On the other hand, in data warehouse, data extracted from transactional system are stored.
03. 80% of the user, can use data warehouse as they need structured data. Structured data is easy to understand and work with. These users use these data to answer specific questions. 10% of the user, do more analysis. Sometimes they need more data then the one available in the data warehouse. Finally, there is user who does deep analysis. They need very large set of data to find out new questions which requires an answer. First two type of user can use data warehouse whereas for the next user they need data lake. However, data lake can support all user types.
04. It is time consuming to change a data warehouse. It requires developer resources. Where is data lake are more adaptable as there is no defined structure or the data is stored in raw format.
05. As all data is available in data lake, user can obtain their result faster than a data warehouse (Campbell 2015).

Step two: Creating a Data Lake Gen2 on Azure

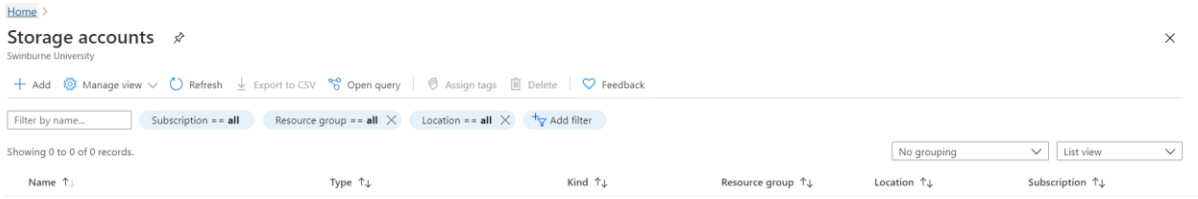

Create storage account

Tables. The cost of your storage account depends on the usage and the options you choose below.
Learn more about Azure storage accounts

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * : Free Trial

Resource group * : 102753762resourcegroup

Create new

Instance details

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. Choose classic deployment model

Storage account name * : 102753762datalake

Location * : (Asia Pacific) Australia East

Performance : Standard / Premium

Account kind : StorageV2 (general purpose v2)

Review + create | < Previous | Next : Networking >

and location does not support large file shares.

Data Lake Storage Gen2

Hierarchical namespace : Disabled / Enabled

NFS v3 : Disabled / Enabled

Sign up is currently required to utilize the NFS v3 feature on a per-subscription basis. Sign up for NFS v3

Create storage account

Validation passed

Basics | Networking | Data protection | Advanced | Tags | Review + create

Basics

Subscription	Free Trial
Resource group	102753762resourcegroup
Location	Australia East
Storage account name	102753762datalake
Deployment model	Resource manager
Account kind	StorageV2 (general purpose v2)
Replication	Read-access geo-redundant storage (RA-GRS)
Performance	Standard
Access tier (default)	Hot

Networking

Connectivity method	Public endpoint (all networks)
Default routing tier	Microsoft network routing (default)

Create | < Previous | Next > | Download a template for au

Delete | Cancel | Redeploy | Refresh

We'd love your feedback! →

Your deployment is complete

Deployment name: Microsoft.StorageAccount-20200819183749
Subscription: Free Trial
Resource group: 102753762resourcegroup

Start time: 8/19/2020, 6:41:06 PM
Correlation ID: 0e61168f-95aa-4a28-a028-a54bd1475953

Deployment details (Download)

Resource	Type	Status	Operation details
102753762datalake	Microsoft.Storage/storageAccoun...	OK	Operation details

Next steps

Go to resource

Step Three: Set up permissions for the managed identity on the Data Lake Storage Gen2

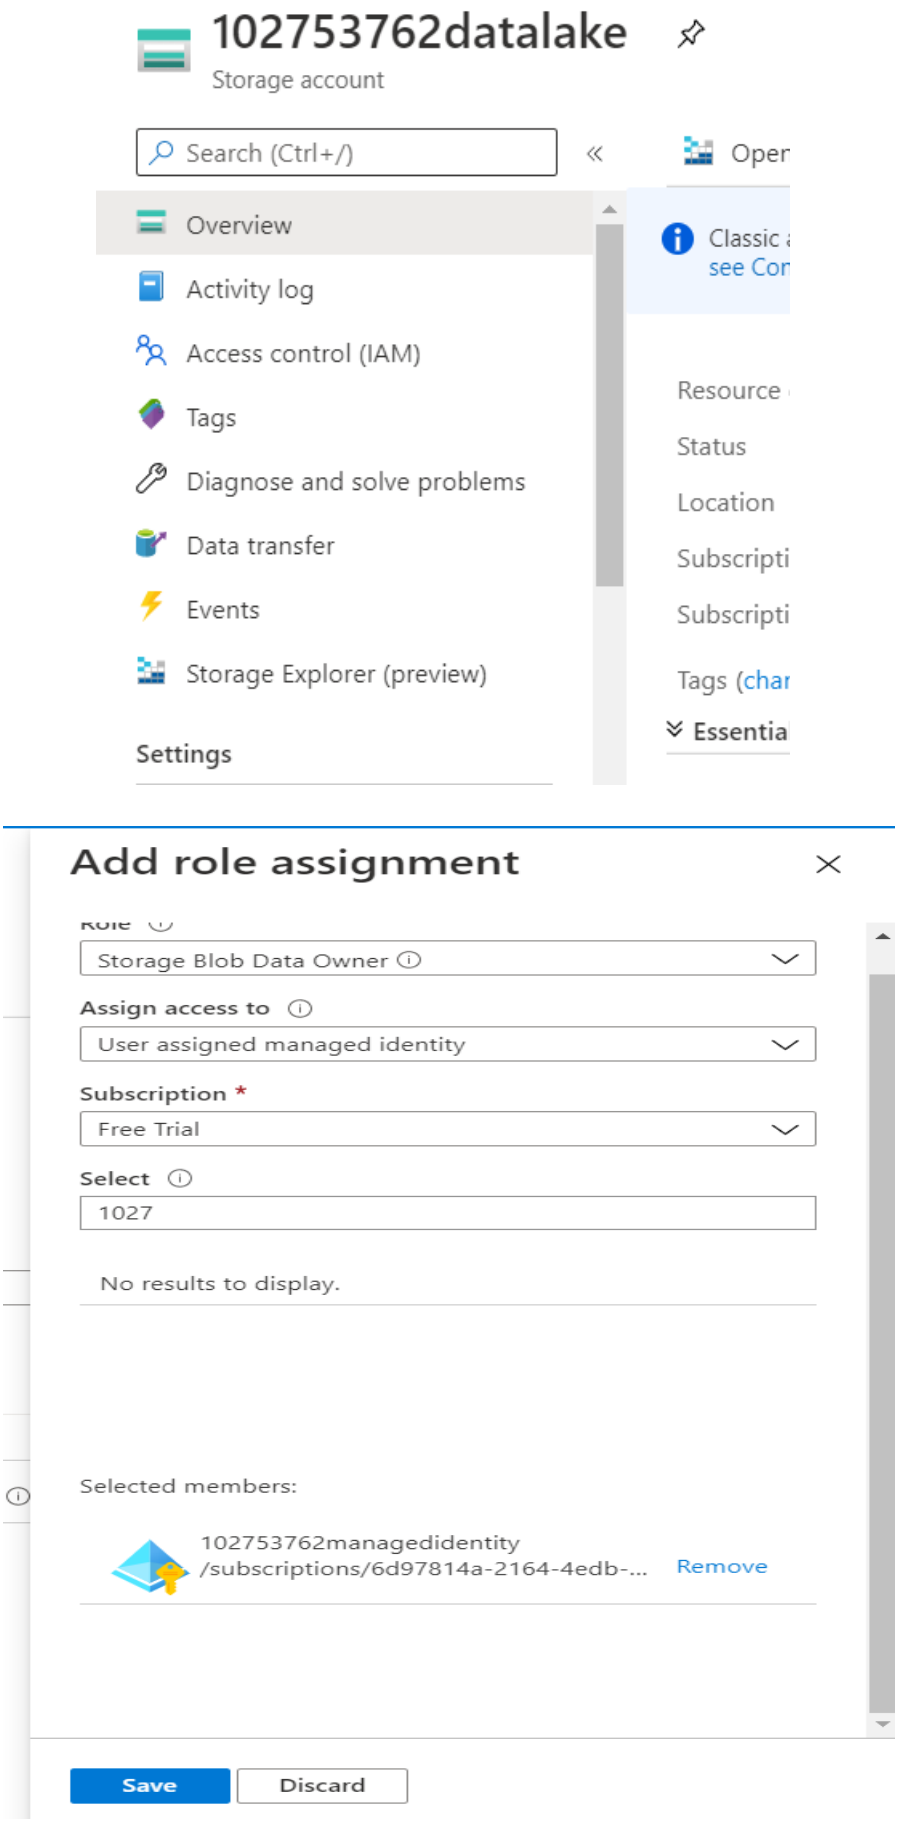

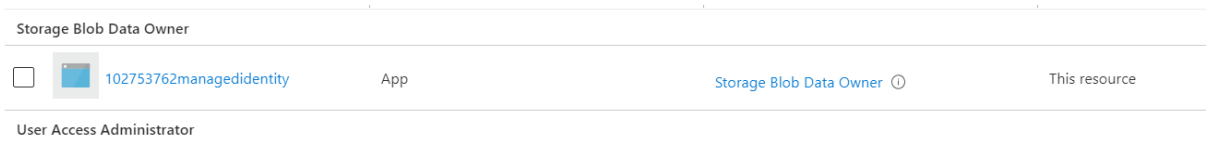

Step Four: Create an SQL Database

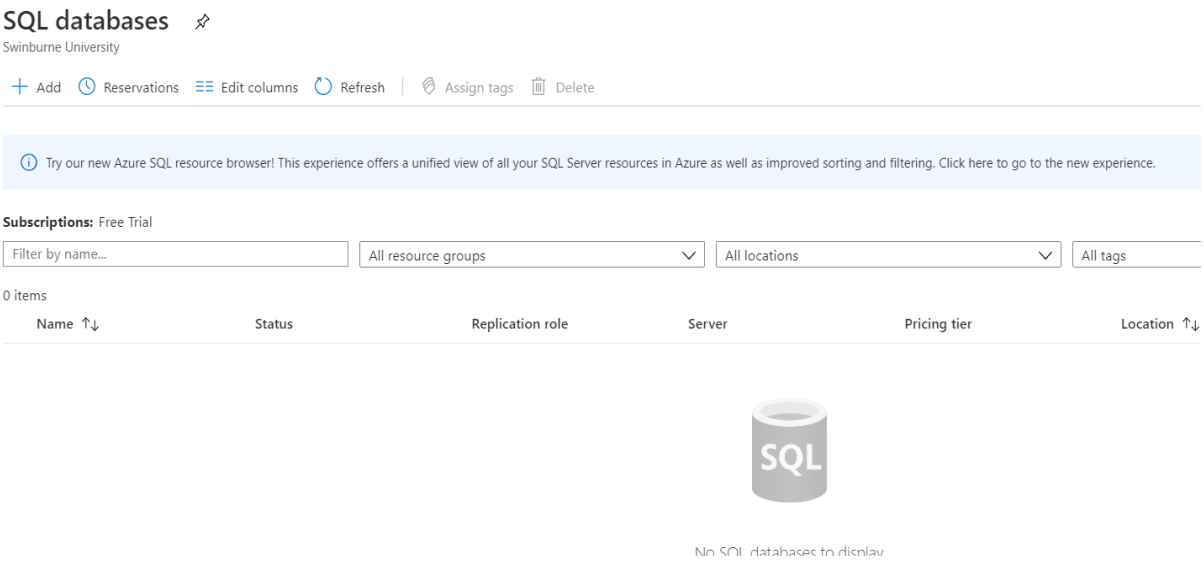

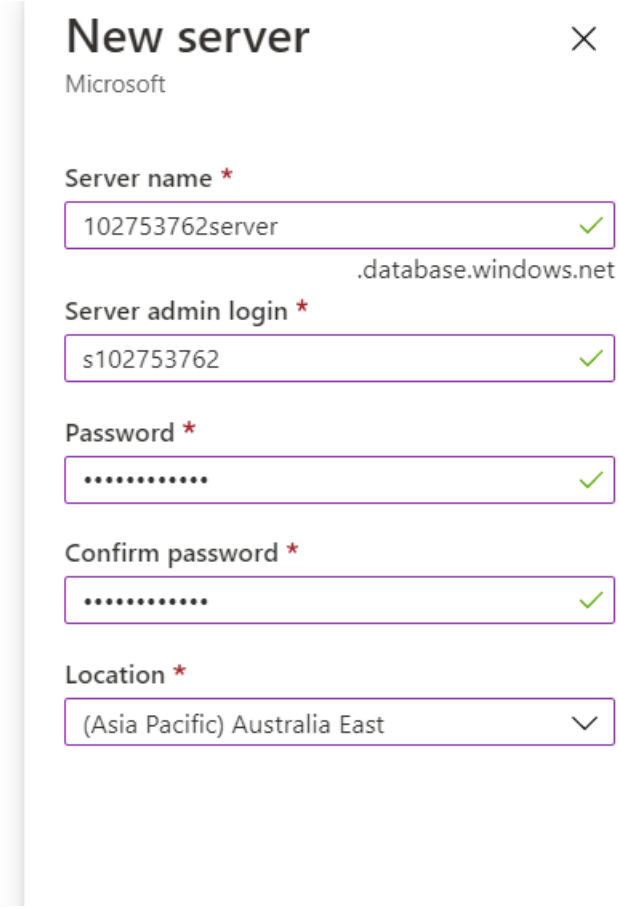

Home > SQL databases >

Create SQL Database

Microsoft

Product details

SQL database
by Microsoft
Terms of use | Privacy policy

Estimated cost per month
22.66 AUD
View pricing details

Terms

By clicking "Create", I (a) agree to the legal terms and privacy statement(s) assc
frequency as my Azure subscription; and (c) agree that Microsoft may share my
party offerings. For additional details see Azure Marketplace Terms.

Basics

Subscription	Free Trial
Resource group	102753762resourcegroup
Region	Australia East
Database name	102753762rdb
Server	(new) 102753762server
Compute + storage	Standard S0: 10 DTUs, 250 GB stora

Create | < Previous | Download a template for automation

Your deployment is complete

Deployment name: Microsoft.SQLDatabase.newDatabaseNewServ...
Subscription: Free Trial
Resource group: 102753762resourcegroup

Start time: 8/19/2020, 6:52:04 PM
Correlation ID: e2679cda-3d08-4a20-95db-43f07cffb417

Deployment details (Download)

Resource	Type	Status	Operation details
102753762server/102753762rdb	Microsoft.Sql/servers/databases	Created	Operation details
102753762server	Microsoft.Sql/servers	OK	Operation details
102753762server	Microsoft.Sql/servers	Created	Operation details

Next steps

Go to resource

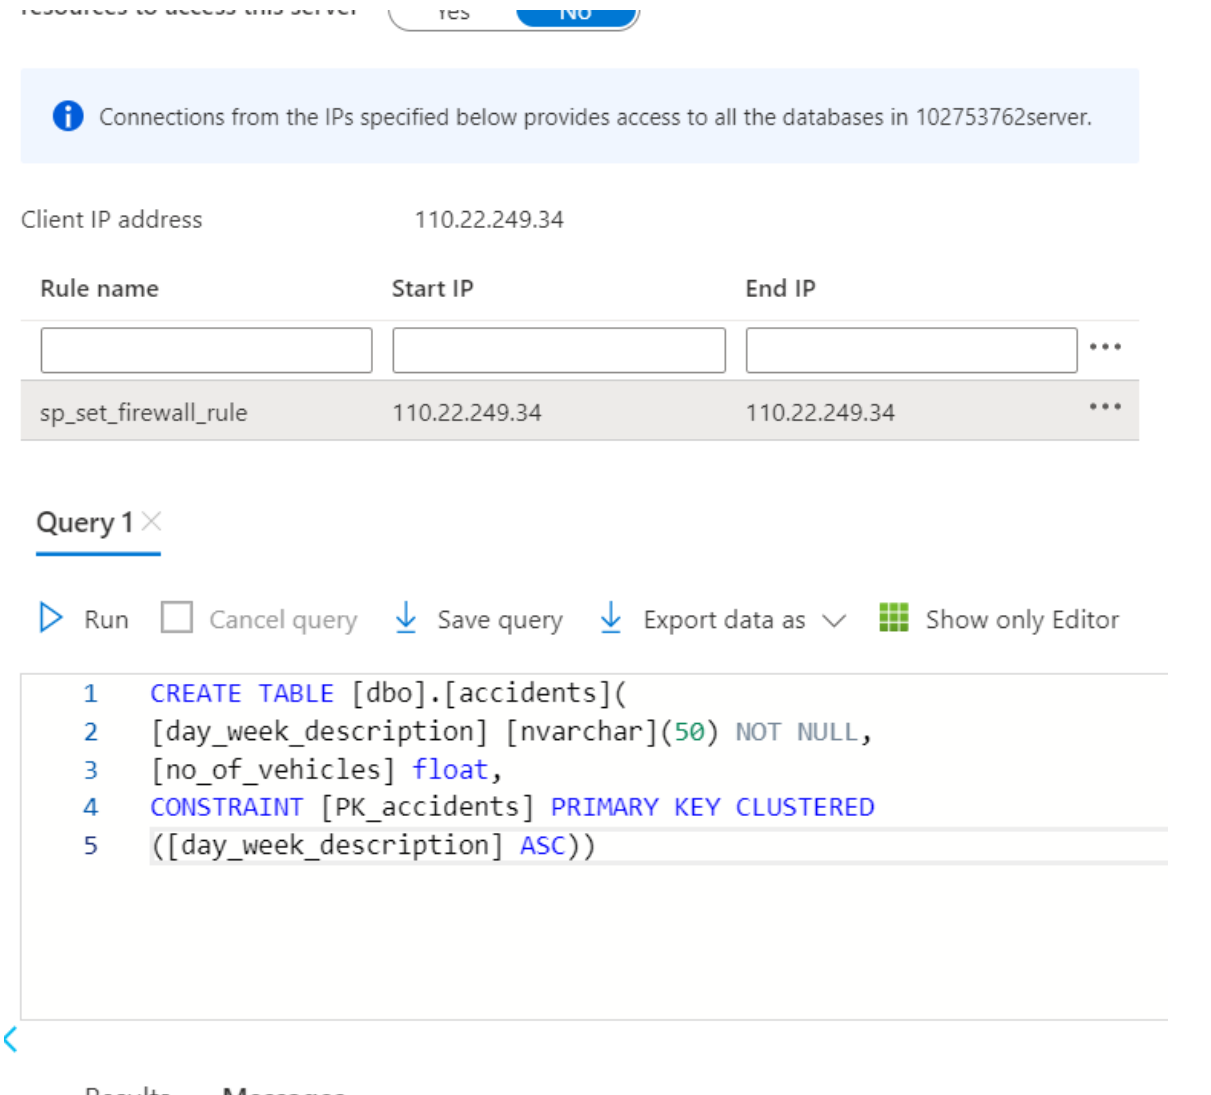

Step Five: Create a HDInsight Cluster

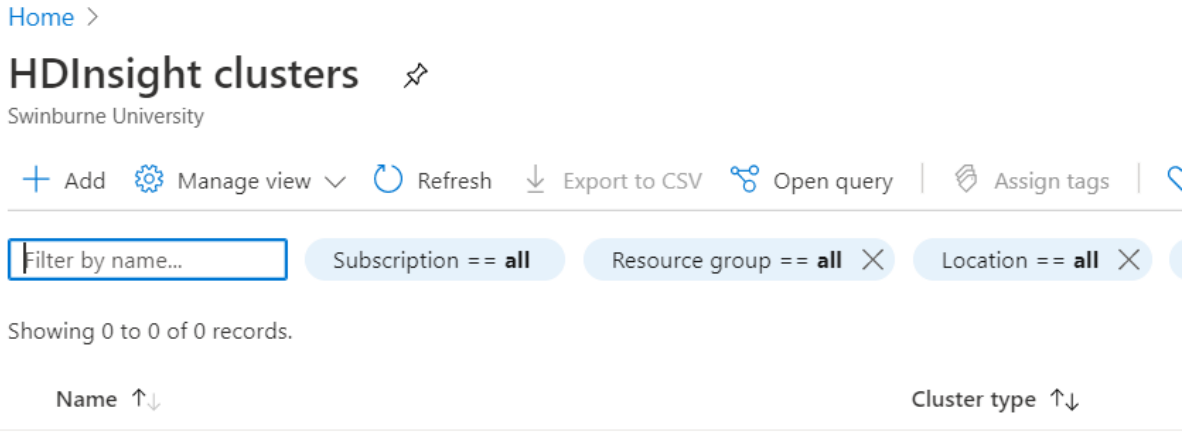

Create HDInsight cluster

Cluster name * s102753762cluster

Region * Australia East

Cluster type * **Hadoop**
Change

Version * Hadoop 2.7.3 (HDI 3.6)

Cluster credentials

Enter new credentials that will be used to administer or access the cluster.

Cluster login username * admin

Cluster login password * ••••••••••••

Confirm cluster login password * ••••••••••••

Secure Shell (SSH) username * sshuser

Use cluster login password for SSH ☑

Review + create « Previous Next: Storage »

Create HDInsight cluster

Basics Storage Security + networking Configuration + pricing Tags Review + create

Select or create storage accounts that will be used for the cluster's logs, job input, and job output. Configure the cluster's access to these accounts, if needed.

Primary storage

Select or create a storage account that will be the default location for cluster logs and other output.

Primary storage type * Azure Data Lake Storage Gen2

Primary storage account * 102753762datalake

Filesystem * 102753762filesystem

Identity

Select a user-assigned managed identity to represent the cluster for Azure Data Lake Gen2 Storage account access. Only identities with access to the selected storage account are listed. Assign the managed identity to the 'Storage Blob Data Owner' role on the storage account. Learn more

User-assigned managed identity * 102753762managedidentity

Additional Azure Storage

Review + create « Previous Next: Security + networking »

Create HDInsight cluster

TLS

Select the minimum TLS version supported for your cluster. Learn more

Minimum TLS version ⓘ 1.2

Network settings

Connect this cluster to a virtual network. Learn more

Virtual network ⓘ

Encryption at rest

Configure disk encryption settings. Learn more

☐ Provide your own key from key vault ⓘ

Identity

Select a user-assigned service identity to represent your cluster for enterprise security package or disk encryption. Learn more

User-assigned managed identity ⓘ 102753762managedidentity

Review + create | « Previous | Next: Configuration + pricing »

Home > HDInsight clusters >

Create HDInsight cluster

✓ Validation succeeded.

Basics | Storage | Security + networking | Configuration + pricing | Tags | Review + create

Hadoop 2.7.3 (HDI 3.6)

5.72 AUD Total estimated cost/hour

This estimate does not include subscription discounts or costs related to s networking, or data transfer.

Basics

Subscription	Free Trial
Resource group	102753762resourcegroup
Region	Australia East
Cluster name	(new) s102753762cluster
Cluster type	Hadoop 2.7.3 (HDI 3.6)
Cluster login username	admin
Secure Shell (SSH) username	sshuser
Use cluster login password for SSH	Enabled

Security + networking

Create | « Previous | Next | Download a template for automation

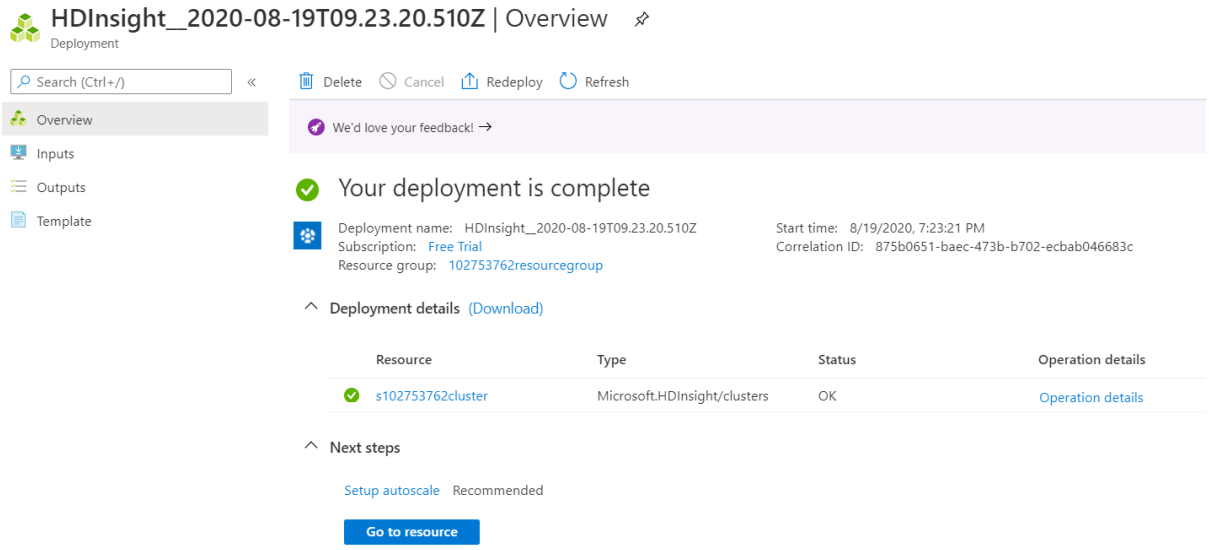

Step Six: Upload the data and staging code

Question 2: What is the script going to do if you run it?

Answer: The script will create a table and store it in textfile format in “data” directory.

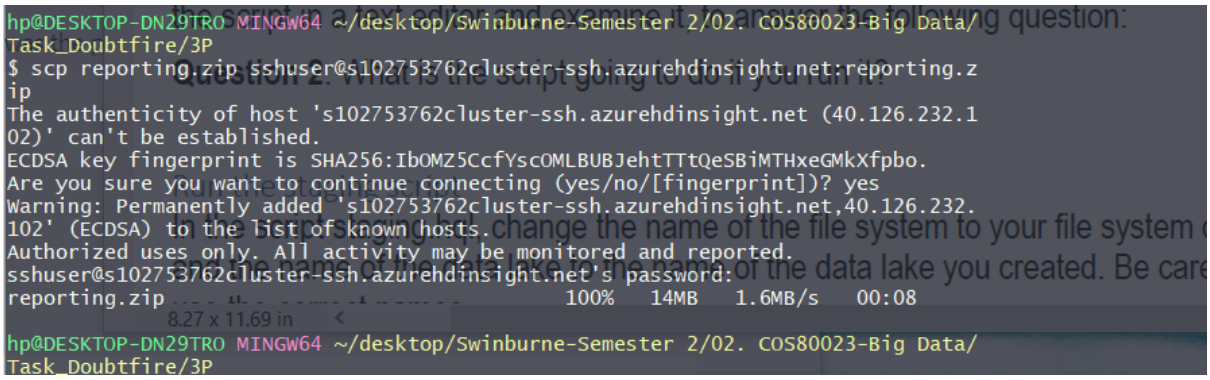

```
hp@DESKTOP-DN29TRO MINGW64 ~/desktop/Swinburne-Semester 2/02. COS80023-Big Data/Task_Doubtfire/3P
$ ssh sshuser@s102753762cluster-ssh.azurehdinsight.net
Authorized uses only. All activity may be monitored and reported.
sshuser@s102753762cluster-ssh.azurehdinsight.net's password:
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.15.0-1091-azure x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 * Are you ready for Kubernetes 1.19? It's nearly here! Try RC3 with
   sudo snap install microk8s --channel=1.19/candidate --classic

   https://microk8s.io/ has docs and details.

0 packages can be updated.
0 updates are security updates.

New release '18.04.5 LTS' available.
Run 'do-release-upgrade' to upgrade to it.


*** /dev/sda1 will be checked for errors at next reboot ***


Welcome to HDInsight.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

sshuser@hn0-s10275:~$
```

```
sshuser@hn0-s10275:~$ unzip reporting.zip
Archive:  reporting.zip
  inflating: reporting.csv
  inflating: staging.hql
sshuser@hn0-s10275:~$ ls
reporting.csv  reporting.zip  staging.hql
sshuser@hn0-s10275:~$
```

```
sshuser@hn0-s10275:~$ hadoop fs -D "fs.azure.createRemoteFileSystemDuringInitialization=true" -ls a
bfs://102753762filesystem@102753762datalake.dfs.core.windows.net/
Found 19 items
-rw-r-----   1 sshuser sshuser          0 2020-08-19 10:42 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/HDInsight_TestAccessiblityBlobName
drwxr-xr-x   - sshuser sshuser          0 2020-08-19 10:55 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/HdiSamples
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/ams
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/amshbase
drwxrwx-wt   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/app-logs
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/apps
drwxr-x--x   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/atshistory
drwxr-xr-x   - sshuser sshuser          0 2020-08-19 10:55 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/custom-scriptaction-logs
drwxr-xr-x   - sshuser sshuser          0 2020-08-19 10:53 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/example
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/hbase
drwxr-x--x   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/hdp
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/hive
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/mapred
drwxrwx-wt   - sshuser sshuser          0 2020-08-19 10:53 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/mapreducestaging
drwxrwx-wt   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/mr-history
drwxrwx-wt   - sshuser sshuser          0 2020-08-19 10:53 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/tezstaging
drwxr-x---   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/tmp
drwxrwx-wt   - sshuser sshuser          0 2020-08-19 10:43 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/user
drwxr-xr-x   - sshuser sshuser          0 2020-08-19 10:52 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/yarn
sshuser@hn0-s10275:~$
```

```
sshuser@hn0-s10275:~$ hdfs dfs -mkdir -p abfs://102753762filesystem@102753762datalake.dfs.core.wind
ows.net/accidents/data
sshuser@hn0-s10275:~$ hdfs dfs -mkdir -p abfs://102753762filesystem@102753762datalake.dfs.core.wind
ows.net/accidents/script
sshuser@hn0-s10275:~$ hdfs dfs -put "reporting.csv" abfs://102753762filesystem@102753762datalake.df
s.core.windows.net/accidents/data/
```

```
sshuser@hn0-s10275:~$ hdfs dfs -ls abfs://102753762filesystem@102753762datalake.dfs.core.windows.n
et/accidents/data

Found 1 items
-rw-r--r--   1 sshuser sshuser  105185364 2020-08-19 11:44 abfs://102753762filesystem@102753762data
lake.dfs.core.windows.net/accidents/data/reporting.csv
sshuser@hn0-s10275:~$
sshuser@hn0-s10275:~$ hdfs dfs -put "staging.hql" abfs://102753762filesystem@102753762datalake.dfs.
core.windows.net/accidents/script/
sshuser@hn0-s10275:~$
```

Step Seven: Transform the Data

```
sshuser@hn0-s10275:~$ beeline -u 'jdbc:hive2://localhost:10001/;transportMode=ht
tp' -f staging.hql
Connecting to jdbc:hive2://localhost:10001/;transportMode=http
Connected to: Apache Hive (version 1.2.1000.2.6.5.3027-5)
Driver: Hive JDBC (version 1.2.1000.2.6.5.3027-5)
Transaction isolation: TRANSACTION_REPEATABLE_READ
0: jdbc:hive2://localhost:10001/> DROP TABLE accidents_raw;
No rows affected (0.58 seconds)
0: jdbc:hive2://localhost:10001/> -- Creates an external table over the csv file

0: jdbc:hive2://localhost:10001/> CREATE EXTERNAL TABLE accidents_raw (ACCIDENT_
NO string,
0: jdbc:hive2://localhost:10001/>     ACCIDENTDATE string,
0: jdbc:hive2://localhost:10001/>     ACCIDENTTIME string,
0: jdbc:hive2://localhost:10001/>     ACCIDENT_TYPE string,
0: jdbc:hive2://localhost:10001/>     Accident_Type_Desc string,
0: jdbc:hive2://localhost:10001/>     DAY_OF_WEEK string,
0: jdbc:hive2://localhost:10001/>     Day_Week_Description string,
0: jdbc:hive2://localhost:10001/>     DCA_CODE string,
0: jdbc:hive2://localhost:10001/>     DCA_Description string,
0: jdbc:hive2://localhost:10001/>     DIRECTORY string,
0: jdbc:hive2://localhost:10001/>     EDITION string,
0: jdbc:hive2://localhost:10001/>     PAGE string,
0: jdbc:hive2://localhost:10001/>     GRID_REFERENCE_X string,
0: jdbc:hive2://localhost:10001/>     GRID_REFERENCE_Y string,
0: jdbc:hive2://localhost:10001/>     LIGHT_CONDITION string,
0: jdbc:hive2://localhost:10001/>     Light_Condition_Desc string,
0: jdbc:hive2://localhost:10001/>     NODE_ID string,
0: jdbc:hive2://localhost:10001/>     NO_OF_VEHICLES float,
0: jdbc:hive2://localhost:10001/>     NO_PERSONS float,
0: jdbc:hive2://localhost:10001/>     NO_PERSONS_INJ_2 string,
0: jdbc:hive2://localhost:10001/>     NO_PERSONS_INJ_3 string,
0: jdbc:hive2://localhost:10001/>     NO_PERSONS_KILLED float,
0: jdbc:hive2://localhost:10001/>     NO_PERSONS_NOT_INJ float,
0: jdbc:hive2://localhost:10001/>     POLICE_ATTEND float,
0: jdbc:hive2://localhost:10001/>     ROAD_GEOMETRY string,
0: jdbc:hive2://localhost:10001/>     Road_Geometry_Desc string,
0: jdbc:hive2://localhost:10001/>     SEVERITY string,
0: jdbc:hive2://localhost:10001/>     SPEED_ZONE float)
0: jdbc:hive2://localhost:10001/> -- The following lines describe the format and
 location of the file
0: jdbc:hive2://localhost:10001/> ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
0: jdbc:hive2://localhost:10001/> LINES TERMINATED BY '\n'
0: jdbc:hive2://localhost:10001/> STORED AS TEXTFILE
0: jdbc:hive2://localhost:10001/> LOCATION 'abfs://102753762filesystem@102753762
datalake.dfs.core.windows.net/accidents/data';
No rows affected (0.775 seconds)
0: jdbc:hive2://localhost:10001/>
0: jdbc:hive2://localhost:10001/> -- Drop the accidents_in_hive table if it exis
ts
0: jdbc:hive2://localhost:10001/> DROP TABLE accidents_in_hive;
No rows affected (0.455 seconds)
0: jdbc:hive2://localhost:10001/> -- Create the accidents_in_hive table and popu
late it with data
```

```
sshuser@hn0-s10275:~$ beeline -u 'jdbc:hive2://localhost:10001/;transportMode=http'
Connecting to jdbc:hive2://localhost:10001/;transportMode=http
Connected to: Apache Hive (version 1.2.1000.2.6.5.3027-5)
Driver: Hive JDBC (version 1.2.1000.2.6.5.3027-5)
Transaction isolation: TRANSACTION_REPEATABLE_READ
Beeline version 1.2.1000.2.6.5.3027-5 by Apache Hive
0: jdbc:hive2://localhost:10001/> INSERT OVERWRITE DIRECTORY '/accidents/output' ROW FORMAT
DELIMITED FIELDS TERMINATED BY '\t' SELECT regexp_replace(day_week_description, '''', ''), s
um(no_of_vehicles) FROM accidents_in_hive WHERE no_of_vehicles IS NOT NULL GROUP BY day_week
_description;
INFO  : Tez session hasn't been created yet. Opening session
DEBUG : Adding local resource: scheme: "hdfs" host: "mycluster" port: -1 file: "/tmp/hive/hi
ve/_tez_session_dir/17291fad-5a0a-4d2f-8ae8-77f931855d1e/hive-hcatalog-core.jar"
INFO  : Dag name: INSERT OVERWRITE DIRE...day_week_description(Stage-1)
DEBUG : DagInfo: {"context":"Hive","description":"INSERT OVERWRITE DIRECTORY '/accidents/out
out' ROW FORMAT DELIMITED FIELDS TERMINATED BY '\\t' SELECT regexp_replace(day_week_descript
ion, '''', ''), sum(no_of_vehicles) FROM accidents_in_hive WHERE no_of_vehicles IS NOT NULL
GROUP BY day_week_description"}
DEBUG : Setting Tez DAG access for queryId=hive_20200820065159_e6fe1a84-60ad-435f-b6ab-3191d
bd616d5 with viewAclString=*, modifyStr=anonymous,hive
INFO  : Status: Running (Executing on YARN cluster with App id application_1597905285696_000
2)

--------------------------------------------------------------------------------
        VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1 ..........  SUCCEEDED      7          7        0        0       0       0
Reducer 2 ......  SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/02  [==========>>] 100%  ELAPSED TIME: 11.30 s
--------------------------------------------------------------------------------
INFO  : Status: DAG finished successfully in 11.08 seconds
INFO  :
INFO  : Query Execution Summary
INFO  : ----------------------------------------------------------------------------------
INFO  : OPERATION                            DURATION
INFO  : ----------------------------------------------------------------------------------
INFO  : Compile Query                          0.51s
INFO  : Prepare Plan                           5.32s
INFO  : Submit Plan                            0.49s
INFO  : Start DAG                              0.39s
INFO  : Run DAG                               11.08s
INFO  : ----------------------------------------------------------------------------------
INFO  :
INFO  : Task Execution Summary
INFO  : ----------------------------------------------------------------------------------
INFO  :   VERTICES  TOTAL_TASKS  FAILED_ATTEMPTS  KILLED_TASKS  DURATION(ms)  CPU_TIME(ms)
 GC_TIME(ms)  INPUT_RECORDS  OUTPUT_RECORDS
INFO  : ----------------------------------------------------------------------------------
INFO  :      Map 1            7                0             0       5896.00        40,170
      1,572        540,253              49
INFO  :  Reducer 2            1                0             0        497.00         1,070
```

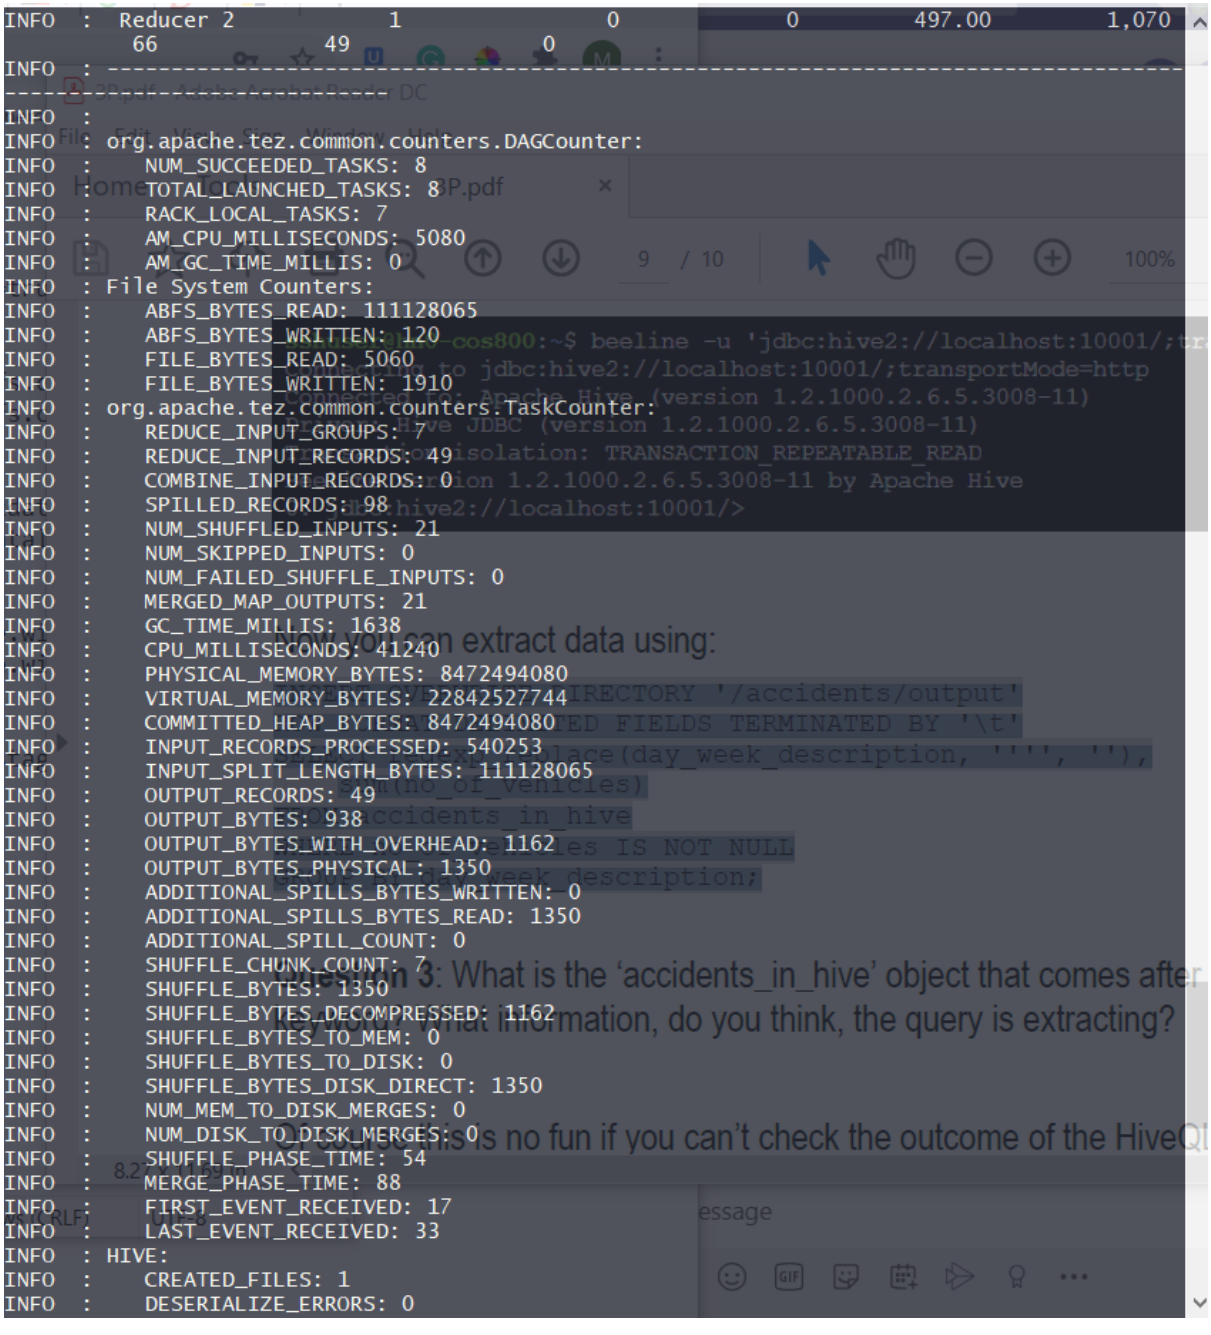
INFO : Reducer 2 1 0 0 497.00 1,070
66 49 0
INFO : org.apache.tez.common.counters.DAGCounter:
INFO : NUM_SUCCEEDED_TASKS: 8
INFO : TOTAL_LAUNCHED_TASKS: 8
INFO : RACK_LOCAL_TASKS: 7
INFO : AM_CPU_MILLISECONDS: 5080
INFO : AM_GC_TIME_MILLIS: 0
INFO : File System Counters:
INFO : ABFS_BYTES_READ: 111128065
INFO : ABFS_BYTES_WRITTEN: 120
INFO : FILE_BYTES_READ: 5060
INFO : FILE_BYTES_WRITTEN: 1910
INFO : org.apache.tez.common.counters.TaskCounter:
INFO : REDUCE_INPUT_GROUPS: 7
INFO : REDUCE_INPUT_RECORDS: 49
INFO : COMBINE_INPUT_RECORDS: 0
INFO : SPILLED_RECORDS: 98
INFO : NUM_SHUFFLED_INPUTS: 21
INFO : NUM_SKIPPED_INPUTS: 0
INFO : NUM_FAILED_SHUFFLE_INPUTS: 0
INFO : MERGED_MAP_OUTPUTS: 21
INFO : GC_TIME_MILLIS: 1638
INFO : CPU_MILLISECONDS: 41240
INFO : PHYSICAL_MEMORY_BYTES: 8472494080
INFO : VIRTUAL_MEMORY_BYTES: 22842527744
INFO : COMMITTED_HEAP_BYTES: 8472494080
INFO : INPUT_RECORDS_PROCESSED: 540253
INFO : INPUT_SPLIT_LENGTH_BYTES: 111128065
INFO : OUTPUT_RECORDS: 49
INFO : OUTPUT_BYTES: 938
INFO : OUTPUT_BYTES_WITH_OVERHEAD: 1162
INFO : OUTPUT_BYTES_PHYSICAL: 1350
INFO : ADDITIONAL_SPILLS_BYTES_WRITTEN: 0
INFO : ADDITIONAL_SPILLS_BYTES_READ: 1350
INFO : ADDITIONAL_SPILL_COUNT: 0
INFO : SHUFFLE_CHUNK_COUNT: 7
INFO : SHUFFLE_BYTES: 1350
INFO : SHUFFLE_BYTES_DECOMPRESSED: 1162
INFO : SHUFFLE_BYTES_TO_MEM: 0
INFO : SHUFFLE_BYTES_TO_DISK: 0
INFO : SHUFFLE_BYTES_DISK_DIRECT: 1350
INFO : NUM_MEM_TO_DISK_MERGES: 0
INFO : NUM_DISK_TO_DISK_MERGES: 0
INFO : SHUFFLE_PHASE_TIME: 54
INFO : MERGE_PHASE_TIME: 88
INFO : FIRST_EVENT_RECEIVED: 17
INFO : LAST_EVENT_RECEIVED: 33
INFO : HIVE:
INFO : CREATED_FILES: 1
INFO : DESERIALIZE_ERRORS: 0

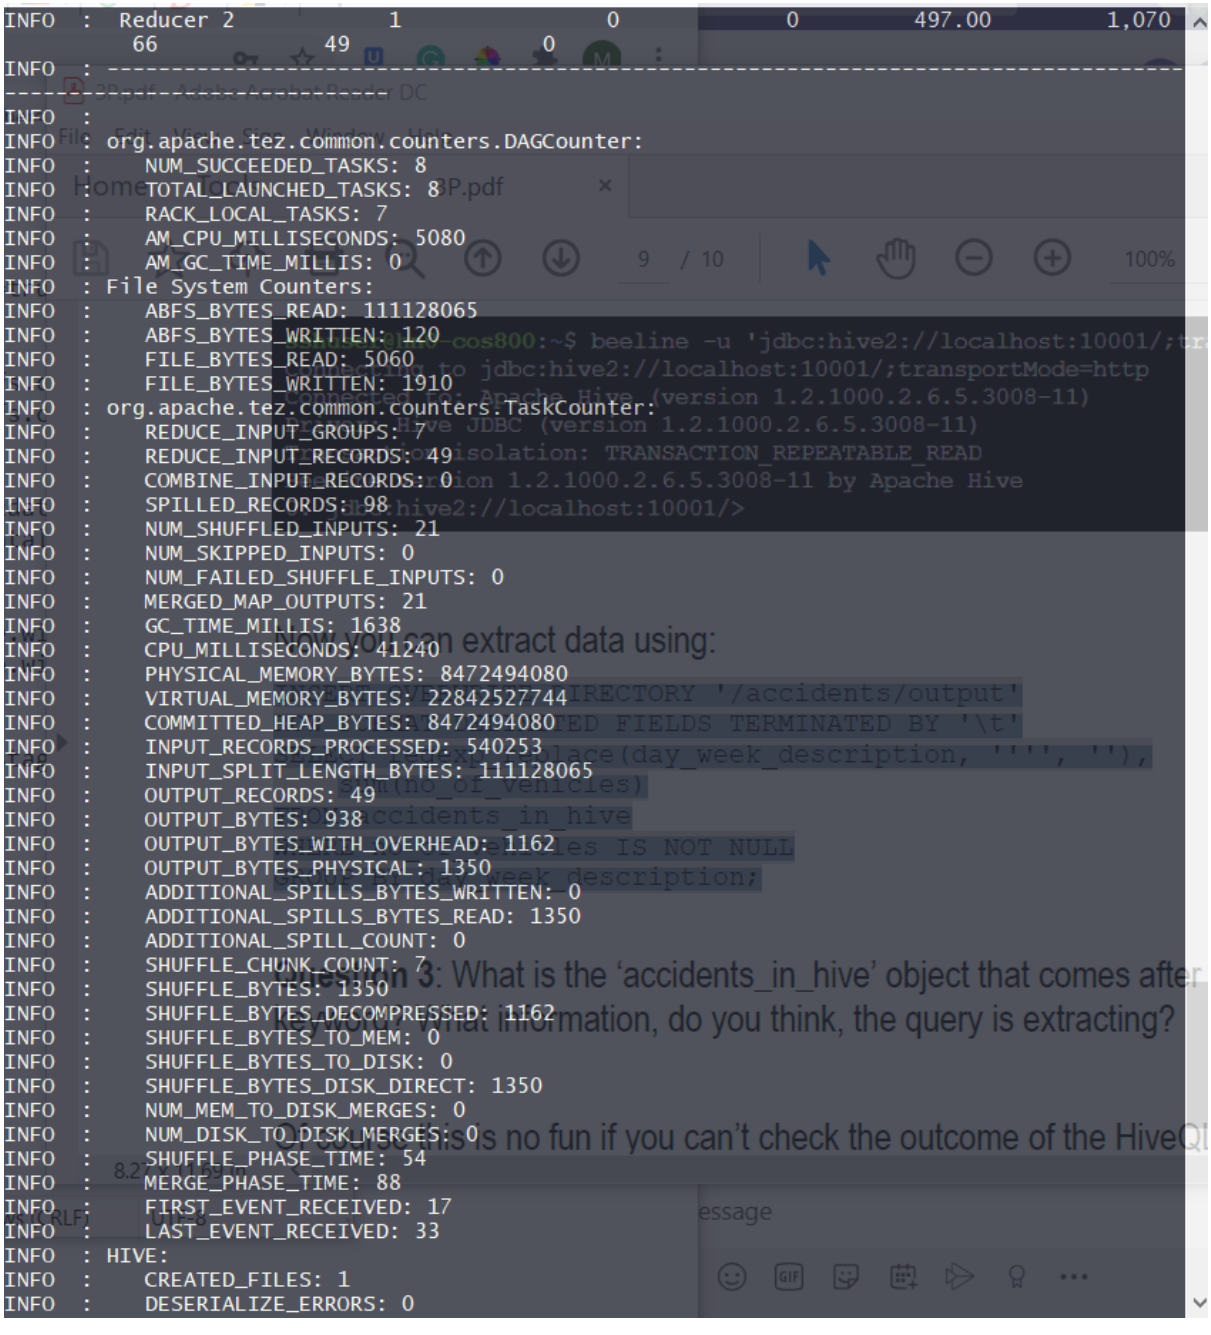
INFO : Reducer 2 1 0 0 497.00 1,070
66 49 0
INFO : org.apache.tez.common.counters.DAGCounter:
INFO : NUM_SUCCEEDED_TASKS: 8
INFO : TOTAL_LAUNCHED_TASKS: 8
INFO : RACK_LOCAL_TASKS: 7
INFO : AM_CPU_MILLISECONDS: 5080
INFO : AM_GC_TIME_MILLIS: 0
INFO : File System Counters:
INFO : ABFS_BYTES_READ: 111128065
INFO : ABFS_BYTES_WRITTEN: 120
INFO : FILE_BYTES_READ: 5060
INFO : FILE_BYTES_WRITTEN: 1910
INFO : org.apache.tez.common.counters.TaskCounter:
INFO : REDUCE_INPUT_GROUPS: 7
INFO : REDUCE_INPUT_RECORDS: 49
INFO : COMBINE_INPUT_RECORDS: 0
INFO : SPILLED_RECORDS: 98
INFO : NUM_SHUFFLED_INPUTS: 21
INFO : NUM_SKIPPED_INPUTS: 0
INFO : NUM_FAILED_SHUFFLE_INPUTS: 0
INFO : MERGED_MAP_OUTPUTS: 21
INFO : GC_TIME_MILLIS: 1638
INFO : CPU_MILLISECONDS: 41240
INFO : PHYSICAL_MEMORY_BYTES: 8472494080
INFO : VIRTUAL_MEMORY_BYTES: 22842527744
INFO : COMMITTED_HEAP_BYTES: 8472494080
INFO : INPUT_RECORDS_PROCESSED: 540253
INFO : INPUT_SPLIT_LENGTH_BYTES: 111128065
INFO : OUTPUT_RECORDS: 49
INFO : OUTPUT_BYTES: 938
INFO : OUTPUT_BYTES_WITH_OVERHEAD: 1162
INFO : OUTPUT_BYTES_PHYSICAL: 1350
INFO : ADDITIONAL_SPILLS_BYTES_WRITTEN: 0
INFO : ADDITIONAL_SPILLS_BYTES_READ: 1350
INFO : ADDITIONAL_SPILL_COUNT: 0
INFO : SHUFFLE_CHUNK_COUNT: 7
INFO : SHUFFLE_BYTES: 1350
INFO : SHUFFLE_BYTES_DECOMPRESSED: 1162
INFO : SHUFFLE_BYTES_TO_MEM: 0
INFO : SHUFFLE_BYTES_TO_DISK: 0
INFO : SHUFFLE_BYTES_DISK_DIRECT: 1350
INFO : NUM_MEM_TO_DISK_MERGES: 0
INFO : NUM_DISK_TO_DISK_MERGES: 0
INFO : SHUFFLE_PHASE_TIME: 54
INFO : MERGE_PHASE_TIME: 88
INFO : FIRST_EVENT_RECEIVED: 17
INFO : LAST_EVENT_RECEIVED: 33
INFO : HIVE:
INFO : CREATED_FILES: 1
INFO : DESERIALIZE_ERRORS: 0

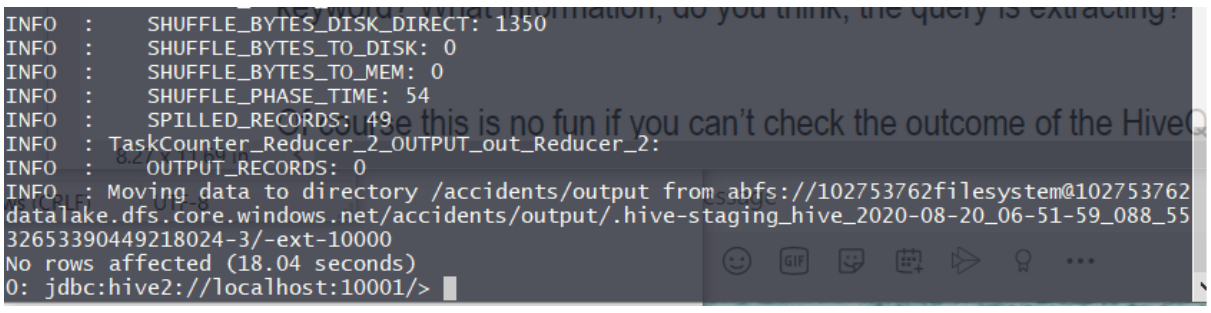
INFO : SHUFFLE_BYTES_DISK_DIRECT: 1350
INFO : SHUFFLE_BYTES_TO_DISK: 0
INFO : SHUFFLE_BYTES_TO_MEM: 0
INFO : SHUFFLE_PHASE_TIME: 54
INFO : SPILLED_RECORDS: 49
INFO : TaskCounter_Reducer_2_OUTPUT_out_Reducer_2:
INFO : OUTPUT_RECORDS: 0
INFO : Moving data to directory /accidents/output from abfs://102753762filesystem@102753762datalake.dfs.core.windows.net/accidents/output/.hive-staging_hive_2020-08-20_06-51-59_088_5532653390449218024-3/-ext-10000
No rows affected (18.04 seconds)
0: jdbc:hive2://localhost:10001/>

Question 3: What is the 'accidents_in_hive' object that comes after the 'FROM' keyword? What information, do you think, the query is extracting?

Answer: "accidents_in_hive" is the table created in hive by running the staging.hql script. The query is extracting data from "accident_in_hive" table to show no of vehicle involved in accident as per weekdays.

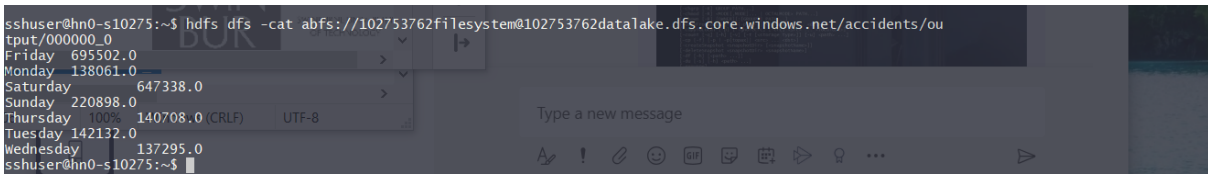

Step Eight: Loading the data

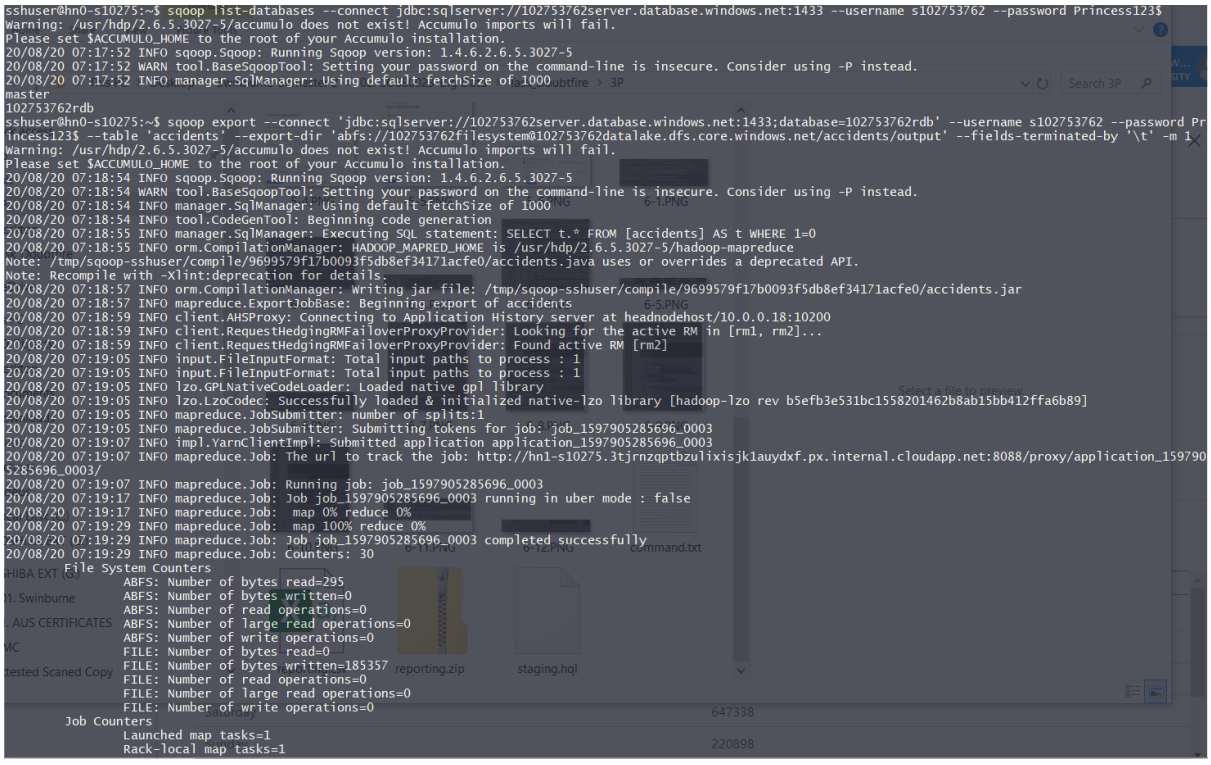

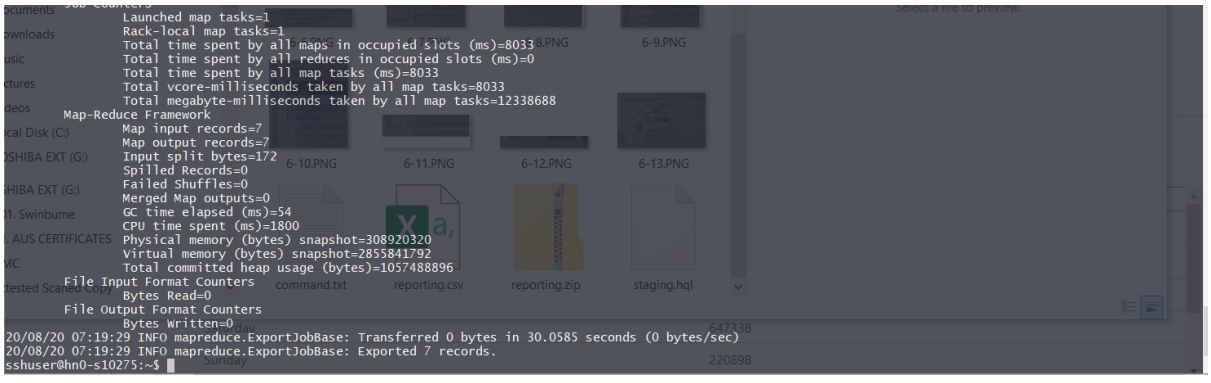

Question 4: On what weekday do the most accidents happen? How many vehicles are involved? Document the answer with a screenshot.

Answer: Most accident happened on Friday and number of vehicles involved were 695502.

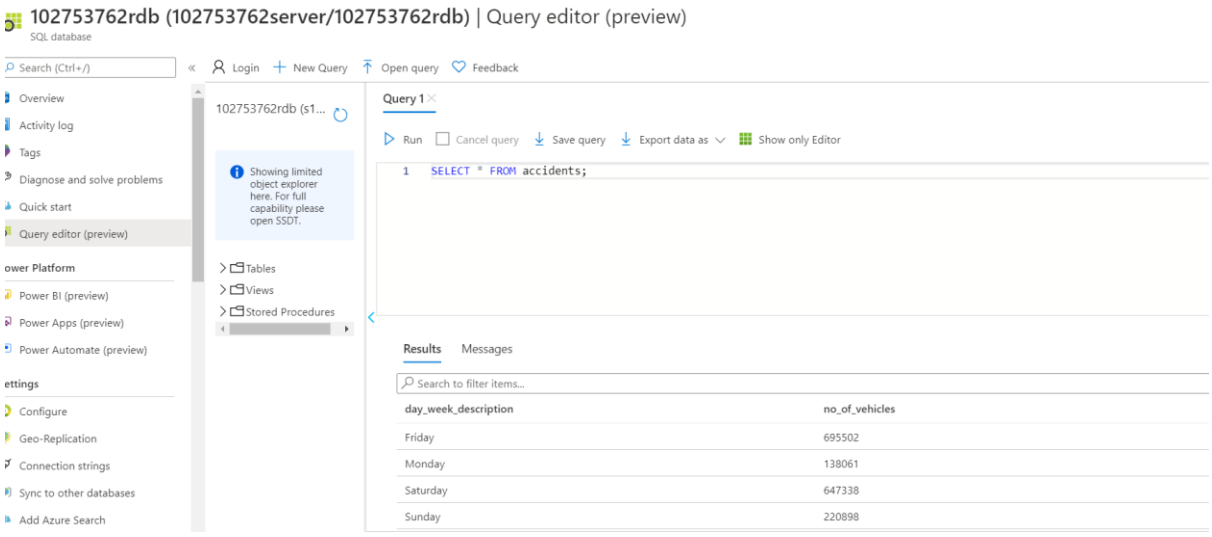

Question 5: What steps and tools were involved in this process? Draw a diagram that shows the tools and files used. You can draw in Powerpoint or by hand and take a photo to document the workflow.

Answer:

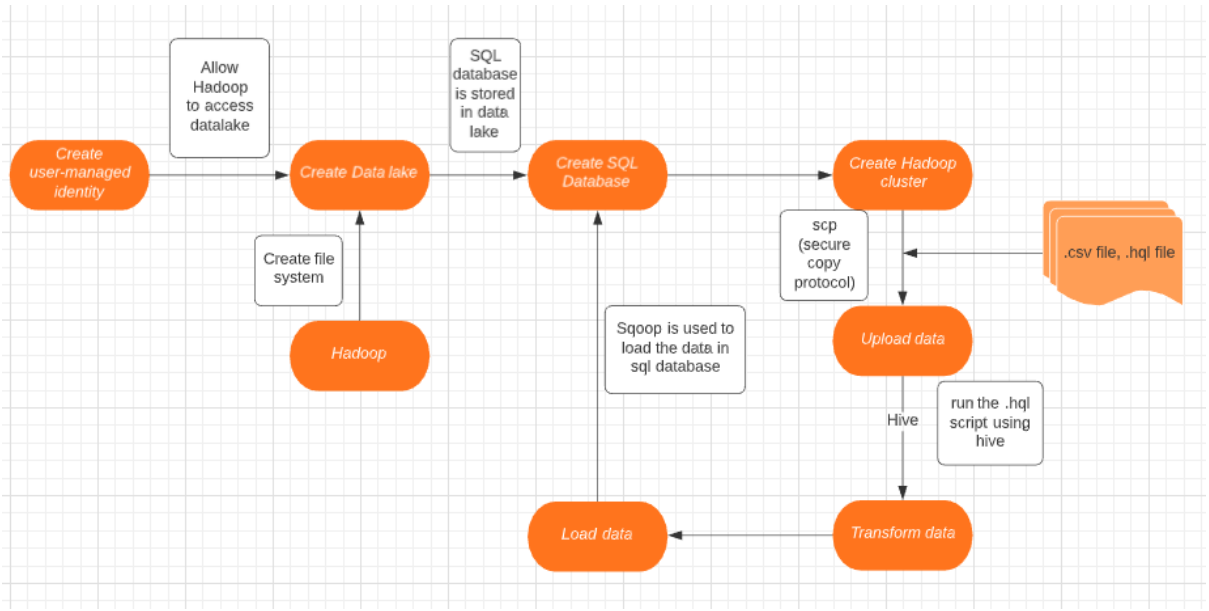

Reference

01. IBM 2020, *What is a Data Warehouse?*, viewed 22 August, 2020, <https://www.ibm.com/cloud/learn/data-warehouse#:~:text=A%20data%20warehouse%20is%20a,AI)%2C%20and%20machine%20learning.&text=Find%20out%20more%20about%20data%20warehouse%20solutions%20from%20IBM.>.
02. Redhat 2020, *What is a data lake?*, viewed 22 August, 2020, <https://www.redhat.com/en/topics/data-storage/what-is-a-data-lake>.

03. Campbell, C 2015, *Top Five Differences between Data Lakes and Data Warehouses*, viewed 22 August, 2020, <https://www.blue-granite.com/blog/bid/402596/top-five-differences-between-data-lakes-and-data-warehouses>.